

# Individual Course Project

## Fall 2025

### MineMind (CLI)

**Objective:** Create a complete, self-contained Minesweeper generator + solver in Python. No GUI, everything runs in a friendly command-line REPL (read–eval–print loop). You’ll implement the game mechanics and a reasoning engine that can solve boards using logic and small exact searches.

### Why This Project?

Learning outcomes (you will be assessed on these):

- Implement classic data structures: Union–Find, priority queues, LRU cache, bitset/bitmask sets, queues/stacks, and use BFS/DFS.
- Model logical constraints and perform deterministic inference (rule-based reasoning).
- Implement a bounded exact search on small components with early pruning, and compute cell-wise probabilities.
- Design a clean CLI/REPL with stable, deterministic behavior and save/load snapshots (JSON).
- Communicate correctness: tests, invariants, complexity analysis, and a concise design write-up.

### What You Will Build

A Python package `minemind` that launches a REPL:

```
python -m minemind new --w 16 --h 16 --mines 40 --seed 1337
```

Then, you play or ask the solver for help using commands like `open 3 7`, `flag 4 9`, `hint`, `step`, `auto`, `--guess`.

**Board sizes:** Beginner 9×9/10 mines, Intermediate 16×16/40, Expert 30×16/99, plus custom.

**First-click safe:** Mines are placed after your first `open x y`, never on that cell or its neighbors.

# Required Features

## CLI commands (REPL)

```
help                                # list commands
new --w W --h H --mines M [--seed S]# start a new game
show [--reveal]                      # print board; --reveal shows mines
(debug/after loss)
open X Y                             # reveal cell at (X,Y)
flag X Y                            # toggle flag
chord X Y                           # on a revealed number: if flags match,
reveal remaining neighbors
hint                                 # print one certain safe/mine move with
explanation
step                                 # apply one deterministic solver step; or
exact small-component step
auto [--guess] [--limit N]          # run solver up to N steps; --guess
allows lowest-risk guesses
prob                                # print coarse ASCII probability heatmap
for unknown cells
frontier                            # summary: #components, sizes, unknowns
per component
save path.json                       # snapshot game state to JSON
load path.json                       # restore snapshot
quit | exit
```

**Coordinates:** zero-based; open 3 7 = column 3, row 7.

## Board printing (ASCII)

- Unknown: · (dot)
- Flag: F
- Revealed number: 1..8
- Revealed zero: space ' '
- Mine (on loss or --reveal): \*

Include column header and row indices for orientation.

## Generator & rules

- Defer mine placement until the first `open`.
- Do not place a mine on the first click or its 8 neighbors.
- Neighbor counts must be correct.
- `open` on a zero performs a flood fill (BFS/DFS) revealing the zero region and its perimeter numbers.
- `chord` on a revealed number reveals all adjacent unknowns only if the number of flagged neighbors equals the number shown.

## Solver requirements

### Frontier & constraints

- Build a frontier consisting of all revealed numbered cells that touch at least one unknown cell.
- For each such cell  $v$  with label  $L$  and unknown neighbors  $N(v)$ :
  - Constraint:  $\sum_{u \in N(v)} x_u = L - \text{flags}(v)$  where  $x_u \in \{0, 1\}$  ( $1 = \text{mine}$ ).
- Represent sets as bitmasks over locally indexed unknowns to allow fast ops.

### Component decomposition (Union–Find)

- Build an intersection graph over constraints: two constraints connect if they share any unknown cell.
- Use union–find to split the frontier into independent components; solve each component separately.

### Deterministic rules (must implement)

1. **Singles**
  - If  $\text{remaining} == 0 \rightarrow$  all cells in scope are **SAFE**.
  - If  $\text{remaining} == |\text{scope}| \rightarrow$  all cells in scope are **MINES**.
2. **Subset rule**
  - For constraints  $(A, a)$  and  $(B, b)$  with  $A \subseteq B$ :
    - If  $a == b \rightarrow B \setminus A$  are **SAFE**.
    - If  $b - a == |B \setminus A| \rightarrow B \setminus A$  are **MINES**.
3. **Scope equality/complements** (merge or simplify identical/complimentary scopes).

Your `hint` must print which rule fired and the involved coordinates.

### Exact enumeration on small components

- For any component with  $k \leq k_{\max}$  unknowns (default `k_max = 20`):
  - Perform backtracking enumeration of satisfying assignments with early pruning.
  - Compute per-cell mine probability = (# solutions with cell=mine) / (total solutions).
  - If `total_solutions == 0`, report an inconsistency (shouldn't occur with correct play).
- Use these probabilities to:
  - Confirm certain safe/mine moves ( $p=0$  or  $p=1$ ) when rules alone don't decide.
  - Support `prob` and `auto --guess`.

## Guess selection (when required)

- Maintain a **min-heap** keyed by `(p_mine, tiebreakers)`; pick the cell with **lowest** estimated mine probability.
- Recommended tiebreakers: prefer central cells, or those that maximize expected information (simple heuristic is fine).

## Caching

- Implement a small LRU cache keyed by a canonical signature of a component (sorted bitmask scopes + remaining vector).
- Invalidate entries when any cell in the component is opened flagged/unflagged.

## Data structures (Required)

- **Union–Find (Disjoint Set)** with path compression + union by rank/size.
- **Priority Queue** (wrap `heappq` cleanly; min-heap).
- **LRU Cache** (hand-rolled or `collections.OrderedDict` based; no third-party lib).
- **Bitmasks** for local sets (ints; use bit ops).
- **Queue/Stack** for flood-fill and undo (if you choose to support undo; optional).
- **Sorting + Binary Search** for deterministic frontier ordering and stable output.

You may use Python **stdlib** only. Do not use SAT/CSP libraries or external DS packages.

## Project Structure

```
/minemind
    __main__.py          # entry: starts REPL
    cli.py               # REPL and command handlers
    render.py            # board ASCII printing
/core
    rng.py               # seeded RNG
    generator.py         # first-click-safe mine placement, neighbor counts
    board.py             # grid state, open/flag/chord, flood fill
    frontier.py          # build frontier, local indexing, component extraction
    dsu.py               # union-find
    rules.py              # singles, subset, scope merges
    solver.py             # exact enumeration, probabilities, auto/step/hint
    lru.py               # LRU cache
    signatures.py         # canonical component signature
    snapshot.py           # save/load JSON snapshot
/tests
    test_generator.py
    test_board.py
    test_rules.py
    test_solver_small.py
    test_frontier.py
```

## Deliverables

1. Code organized as above (or equivalent clarity).
2. README.md
  - o How to install/run.
  - o Command cheatsheet with examples.
  - o Known limitations.
3. DESIGN.md
  - o Diagrams of frontier/components.
  - o Invariants for DSU, board, and solver.
  - o Data flow for hint/step/auto.
4. COMPLEXITY.md
  - o Big-O for core operations (open, flood, build frontier, rule pass, enumeration).
  - o Justify your  $k_{\max}$  choice (e.g., 20).
5. TESTING.md
  - o What each test covers; how to run pytest.
6. At least 15 unit tests hitting: generator properties, flood fill, rules, DSU components, small exact solve, probability sums.
7. Two sample snapshots (.json): one mid-game; one near endgame.
8. Short demo (text transcript or <3-min screen capture) showing: new, open, a rule-based hint, an auto --guess, save/load.

## Grading rubric (100 pts + up to 7 bonus)

- **Functionality** (40 pts)
  - o Generator (first-click safety, counts), play features (open/flag/chord/flood) ... 15
  - o Solver deterministic rules (hint, step) ... 10
  - o Exact enumeration & probabilities; auto --guess ... 15
- **DS/Algorithm Correctness** (25 pts)
  - o DSU partitioning & invariants ... 6
  - o Bitset subset logic & rule soundness ... 8
  - o Enumeration correctness, pruning, probability math ... 11
- **Code Quality** (15 pts)
  - o Modularity, naming, docstrings, invariants, typing ... 15
- **Testing** (10 pts)
  - o  $\geq 15$  tests; good edge cases; deterministic behavior with seeds ... 10
- **Documentation** (10 pts)
  - o README, DESIGN, COMPLEXITY, TESTING completeness & clarity ... 10
- **Bonus** (up to +7 pts)
  - o Monte Carlo fallback for  $k > 20$  components ... +3
  - o 3BV difficulty metric & histogram ... +2
  - o Verbose “proof viewer” in hint (shows coordinates/bitsets used) ... +2

## Acceptance criteria (Allows You to Self-Check Your Solution)

- **First click safety:** In 200 randomized trials, first click and its neighbors are never mines.
- **Counts correct:** Random boards pass property `test_count == #adjacent_mines`.
- **Chord correctness:** Only reveals when flags match the number.
- **Rules fire:** Provide at least 3 reproducible scenarios where Singles and Subset both trigger.
- **Exact solve:** On a component with  $k \leq 10$ , probabilities match brute force within floating-point tolerance.
- **auto --guess:** Works and terminates on complete game or step limit; prints final status.
- **Save/Load:** Round-trips without loss (same board state and clock).
- **Determinism:** With seed fixed, repeated sessions produce the same sequence of boards.

## Implementation tips

- **Bitmasks:** map local unknowns  $0..k-1 \rightarrow \text{int mask}$ ; use `&`, `|`, `^`, `~`, `bit_count()`.
- **Pruning:** maintain `remaining` for each constraint; if `remaining < 0` or `remaining > |scope.unassigned|`, backtrack.
- **Branching order:** choose the variable in most constraints first to reduce search.
- **LRU:** cache only after sorting scopes and normalizing remaining vector.
- **Testing harness:** seed PRNG; generate random boards; assert invariants before/after moves.

## Policies

### Allowed resources

- Python **3.10+** (prefer 3.11).
- Standard library only (e.g., `random`, `heapq`, `collections`, `dataclasses`, `argparse`, `json`).
- You may consult references/tutorials for Minesweeper rules and general DS ideas, but **core solver/DS code must be developed by you**.

### AI/assistance policy

- You may use AI tools **for explanations or debugging ideas**, but **not** to generate the core DS/solver code.
- If you used any external sources for ideas, add a short note in `ACKNOWLEDGMENTS.md`.

## Collaboration

- Discussion of concepts is OK; sharing code is not. Your tests and implementation must be your own.

## Submission

- Submit a zip or repository link that includes: code, tests, snapshots, and all required docs.
- Ensure `python -m minemind new --w 9 --h 9 --mines 10 --seed 42` launches the REPL and `help` works.
- Include a short transcript in README showing a typical session.

## Example Session

```
$ python -m minemind new --w 9 --h 9 --mines 10 --seed 42
New game: 9x9, 10 mines, seed=42
> show
  0 1 2 3 4 5 6 7 8
0 . . . . . . .
...
> open 3 3
(revealed 12 cells)
  0 1 2 3 4 5 6 7 8
0     1 1 1 . . .
1     1     1 . . .
2     1 1 1 . . .
3     1     1 . . .
...
> hint
SAFE: (5,4) — SUBSET: N(4,3) ⊆ N(4,4) and remaining equal → B\A safe
> step
Applied SINGLE at (2,5): remaining=0 → all neighbors safe
> auto --guess --limit 100
Auto: 27 steps; guessed (1,8) with p_mine=0.08; solved in 31 steps.
> save run1.json
Saved.
```

## Glossary

- **Frontier:** set of numbered revealed cells adjacent to unknowns.
- **Component:** independent subset of constraints/unknowns (no overlap with others).
- **k\_max:** max unknowns in a component for exact enumeration.
- **3BV:** board difficulty metric (bonus).